

Course 2 – Week 2 – Quiz

Quiz 1 – Neural Network Training

✓ Congratulations! You passed!

Grade
received 100%

Latest Submission
Grade 100%

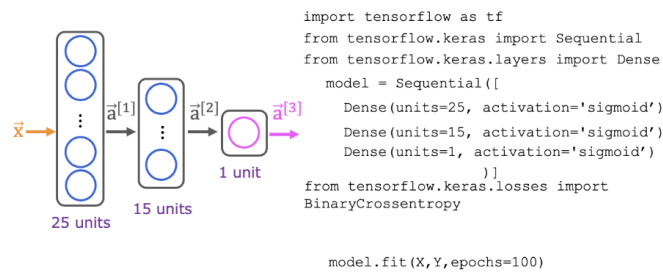
To pass 80% or
higher

Go to next item

1.

1 / 1 point

Train a Neural Network in TensorFlow



Here is some code that you saw in the lecture:

```
...

model.compile(loss=BinaryCrossentropy())

...
```

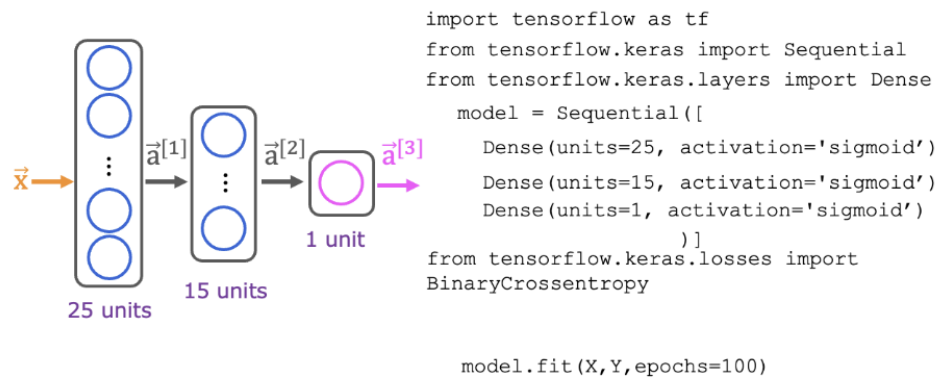
For which type of task would you use the binary cross entropy loss function?

- ☒ binary classification (classification with exactly 2 classes)
- ☐ regression tasks (tasks that predict a number)
- ☐ BinaryCrossentropy() should not be used for any task.
- ☐ A classification task that has 3 or more classes (categories)

✓ Correct

Yes! Binary cross entropy, which we've also referred to as logistic loss, is used for classifying between two classes (two categories).

Train a Neural Network in TensorFlow



Here is code that you saw in the lecture:

```

...

model = Sequential([
    Dense(units=25, activation='sigmoid'),
    Dense(units=15, activation='sigmoid'),
    Dense(units=1, activation='sigmoid')
])

model.compile(loss=BinaryCrossentropy())

model.fit(X,y,epochs=100)

...

```

Which line of code updates the network parameters in order to reduce the cost?

- ☐ `model.compile(loss=BinaryCrossentropy())`
- ☐ `model = Sequential([...])`
- ☒ `model.fit(X,y,epochs=100)`
- ☐ None of the above -- this code does not update the network parameters.

✓ **Correct**


Yes! The third step of model training is to train the model on data in order to minimize the loss (and the cost)

Quiz 2 – Activation Functions

1.

1 / 1 point

Examples of Activation Functions




Which of the following activation functions is the most common choice for the hidden layers of a neural network?

- ☐ Sigmoid
- ☒ ReLU (rectified linear unit)
- ☐ Most hidden layers do not use any activation function
- ☐ Linear


✓ Correct

Yes! A ReLU is most often used because it is faster to train compared to the sigmoid. This is because the ReLU is only flat on one side (the left side) whereas the sigmoid goes flat (horizontal, slope approaching zero) on both sides of the curve.


2.




Binary classification
Sigmoid
 $y = 0/1$



Regression
Linear activation function
 $y = +/-$



Regression
ReLU
 $Y = 0 \text{ or } +$



For the task of predicting housing prices, which activation functions could you choose for the output layer?
Choose the 2 options that apply.

☐ Sigmoid

☒ ReLU

✓ **Correct**

Yes! ReLU outputs values 0 or greater, and housing prices are positive values.

☐ linear

You didn't select all the correct answers

3. True/False? A neural network with many layers but no activation function (in the hidden layers) is not effective; that's why we should instead use the linear activation function in every hidden layer.

1 / 1 point

☐ True

☒ False

✓ **Correct**

Yes! A neural network with many layers but no activation function is not effective. A linear activation is the same as "no activation function".

Quiz 3 – Multiclass Classification

1.

1 / 1 point

Softmax regression (4 possible outputs)

$$\begin{aligned} \times z_1 &= \vec{w}_1 \cdot \vec{x} + b_1 & a_1 &= \frac{e^{z_1}}{e^{z_1} + e^{z_2} + e^{z_3} + e^{z_4}} \\ & & &= P(y = 1|\vec{x}) \quad 0.30 \\ \circ z_2 &= \vec{w}_2 \cdot \vec{x} + b_2 & a_2 &= \frac{e^{z_2}}{e^{z_1} + e^{z_2} + e^{z_3} + e^{z_4}} \\ & & &= P(y = 2|\vec{x}) \quad 0.20 \\ \square z_3 &= \vec{w}_3 \cdot \vec{x} + b_3 & a_3 &= \frac{e^{z_3}}{e^{z_1} + e^{z_2} + e^{z_3} + e^{z_4}} \\ & & &= P(y = 3|\vec{x}) \quad 0.15 \\ \triangle z_4 &= \vec{w}_4 \cdot \vec{x} + b_4 & a_4 &= \frac{e^{z_4}}{e^{z_1} + e^{z_2} + e^{z_3} + e^{z_4}} \\ & & &= P(y = 4|\vec{x}) \quad 0.35 \end{aligned}$$

For a multiclass classification task that has 4 possible outputs, the sum of all the activations adds up to 1. For a multiclass classification task that has 3 possible outputs, the sum of all the activations should add up to

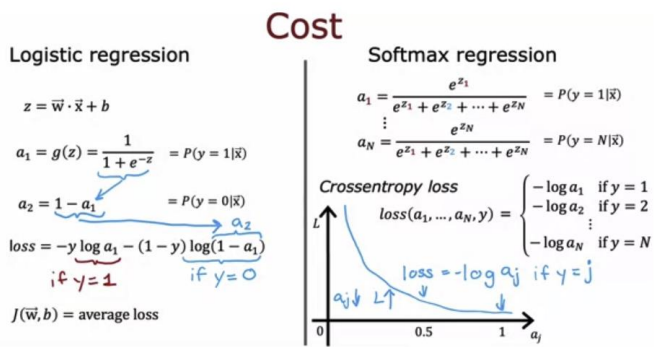
- ☐ More than 1
- ☒ 1
- ☐ It will vary, depending on the input x .
- ☐ Less than 1

✓ Correct

Yes! The sum of all the softmax activations should add up to 1. One way to see this is that if $e^{z_1} = 10, e^{z_2} = 20, e^{z_3} = 30$, then the sum of $a_1 + a_2 + a_3$ is equal to $\frac{e^{z_1} + e^{z_2} + e^{z_3}}{e^{z_1} + e^{z_2} + e^{z_3}}$ which is 1.

2.

1 / 1 point



For multiclass classification, the cross entropy loss is used for training the model. If there are 4 possible classes for the output, and for a particular training example, the true class of the example is class 3 ($y=3$), then what does the cross entropy loss simplify to? [Hint: This loss should get smaller when a_3 gets larger.]

- ☐ $z_3 / (z_1 + z_2 + z_3 + z_4)$
- ☐ z_3
- ☐ $\frac{-\log(a_1) - \log(a_2) - \log(a_3) - \log(a_4)}{4}$
- ☒ $-\log(a_3)$

✓ Correct

Correct. When the true label is 3, then the cross entropy loss for that training example is just the negative of the log of the activation for the third neuron of the softmax. All other terms of the cross entropy loss equation ($-\log(a_1)$, $-\log(a_2)$, and $-\log(a_4)$) are ignored

3.

1 / 1 point

MNIST (more numerically accurate)

```
model    import tensorflow as tf
         from tensorflow.keras import Sequential
         from tensorflow.keras.layers import Dense
         model = Sequential([
             Dense(units=25, activation='relu')
             Dense(units=15, activation='relu')
             Dense(units=10, activation='linear') ])

loss     from tensorflow.keras.losses import
         SparseCategoricalCrossentropy

         model.compile(...,loss=SparseCategoricalCrossentropy(from_logits=True) )

fit      model.fit(X,Y,epochs=100)

predict  logits = model(X)
         f_x = tf.nn.softmax(logits)
```

For multiclass classification, the recommended way to implement softmax regression is to set `from_logits=True` in the loss function, and also to define the model's output layer with...

- ☐ a 'softmax' activation
- ☒ a 'linear' activation

✓ **Correct**

Yes! Set the output as linear, because the loss function handles the calculation of the softmax with a more numerically stable method.

Quiz 4 – Additional Neural Network Concepts

1.

MNIST Adam

1 / 1 point

```
model
model = Sequential([
    tf.keras.layers.Dense(units=25, activation='sigmoid')
    tf.keras.layers.Dense(units=15, activation='sigmoid')
    tf.keras.layers.Dense(units=10, activation='linear')
])

compile
 $\alpha = 10^{-3} = 0.001$ 
model.compile(optimizer=tf.keras.optimizers.Adam(learning_rate=1e-3),
              loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True))

fit
model.fit(X, Y, epochs=100)
```

The Adam optimizer is the recommended optimizer for finding the optimal parameters of the model. How do you use the Adam optimizer in TensorFlow?


- ☒ When calling model.compile, set optimizer=tf.keras.optimizers.Adam(learning_rate=1e-3).
- ☐ The call to model.compile() uses the Adam optimizer by default
- ☐ The Adam optimizer works only with Softmax outputs. So if a neural network has a Softmax output layer, TensorFlow will automatically pick the Adam optimizer.
- ☐ The call to model.compile() will automatically pick the best optimizer, whether it is gradient descent, Adam or something else. So there's no need to pick an optimizer manually.

✓ Correct
Correct. Set the optimizer to Adam.

2.

Convolutional Neural Network

1 / 1 point



The lecture covered a different layer type where each single neuron of the layer does not look at all the values of the input vector that is fed into that layer. What is this name of the layer type discussed in lecture?

- ☐ 1D layer or 2D layer (depending on the input dimension)
- ☒ convolutional layer
- ☐ Image layer
- ☐ A fully connected layer

✓ Correct
Correct. For a convolutional layer, each neuron takes as input a subset of the vector that is fed into that layer.