Congratulations! You passed!

Grade received 100% **Latest Submission Grade** 100% **To pass** 80% or higher

Go to next item

Quiz

Continuous state spaces

Submit your assignment

Due Mar 10, 11:59 PM CET

Try again

Receive grade

To Pass 80% or higher

Your grade

100%

View Feedback

We keep your highest score

Like Dislike Report an issue

1. 1 / 1 point

The Lunar Lander is a continuous state Markov Decision Process (MDP) because:

- ○ The state-action value $Q(s, a)$ function outputs continuous valued numbers
- ○ The state has multiple numbers rather than only a single number (such as position in the x-direction)
- ◉ The state contains numbers such as position and velocity that are continuous valued.
- ○ The reward contains numbers that are continuous valued

Correct
That's right!

2. 1 / 1 point

In the learning algorithm described in the videos, we repeatedly create an artificial training set to which we apply supervised learning where the input $x = (s, a)$ and the target, constructed using Bellman's equations, is y = _____?

- ◉ $y = R(s) + \gamma \max_{a'} Q(s', a')$ where s' is the state you get to after taking action a in state s
- ○ $y = R(s)$
- ○ $y = R(s')$ where s' is the state you get to after taking action a in state s
- ○ $y = \max_{a'} Q(s', a')$ where s' is the state you get to after taking action a in state s

Correct

3. 1 / 1 point

You have reached the final practice quiz of this class! What does that mean? (Please check all the answers, because all of them are correct!)

- ☑ Andrew sends his heartfelt congratulations to you!

Correct

- ☑ You deserve to celebrate!

Correct

- ☑ The DeepLearning.AI and Stanford Online teams would like to give you a round of applause!

Correct

- ☑ What an accomplishment -- you made it!

Correct